

Predictive Analytics & Data Mining on Campus Placements – Team 4

Final Assignment

Team Members:

Prashanth ASOK KUMAR

Annanya CHITRA KANNAN

Levon AVETISYAN

Vasanth MOHAN

Introduction

We are a group of 4, and we have split our task to work on both R-studio and Dataiku. Our data set is regarding campus placement of students at Jain University in Bangalore, India.

Link: <https://www.kaggle.com/benroshan/factors-affecting-campus-placement>

R-Studio Part

We are trying to try to predict using if the student will be placed or not and what will be his salary if placed. We are using the following 3 models for prediction.

- Linear Regression Model
- Knn classification Model
- Naive Bayes Model

Before starting with the prediction, lets see if there are any relation of the data points with status and salary.

Libraries used:

```
library(e1071)
library(ggplot2)
library(tidyr)
library(gmodels)
library(class)
library(dplyr)

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

Importing Data and Data Correction:

Removing the Serial Number Column and replacing the NA in salary column with 0

```

placement<-read.csv("Placement_Data_Full_Class.csv", stringsAsFactors=TRUE)
placement<-placement[-1]
placement[is.na(placement)] <- 0
str(placement)

## 'data.frame':    215 obs. of  14 variables:
## $ gender       : Factor w/ 2 levels "F","M": 2 2 2 2 2 2 1 2 2 2 ...
## $ ssc_p        : num  67 79.3 65 56 85.8 ...
## $ ssc_b        : Factor w/ 2 levels "Central","Others": 2 1 1 1 1 2 2 1
1 1 ...
## $ hsc_p        : num  91 78.3 68 52 73.6 ...
## $ hsc_b        : Factor w/ 2 levels "Central","Others": 2 2 1 1 1 2 2 1
1 1 ...
## $ hsc_s        : Factor w/ 3 levels "Arts","Commerce",...: 2 3 1 3 2 3 2
3 2 2 ...
## $ degree_p     : num  58 77.5 64 52 73.3 ...
## $ degree_t     : Factor w/ 3 levels "Comm&Mgmt","Others",...: 3 3 1 3 1 3
1 3 1 1 ...
## $ workex       : Factor w/ 2 levels "No","Yes": 1 2 1 1 1 2 1 2 1 1 ...
## $ etest_p      : num  55 86.5 75 66 96.8 ...
## $ specialisation: Factor w/ 2 levels "Mkt&Fin","Mkt&HR": 2 1 1 2 1 1 1 1 1
1 1 ...
## $ mba_p        : num  58.8 66.3 57.8 59.4 55.5 ...
## $ status       : Factor w/ 2 levels "Not Placed","Placed": 2 2 2 1 2 1 1
2 2 1 ...
## $ salary       : num  270000 200000 250000 0 425000 0 0 252000 231000 0
...

```


Relation of data points with salary and status:

Relationship with status: The following have linearly positive relation with status:

```

placement%>%gather(placement,nums,c("ssc_p","hsc_p","degree_p" ))%>%
  ggplot(aes(x=status,y=nums))+geom_boxplot() + facet_wrap(~ placement,
scales = "free_y")

```



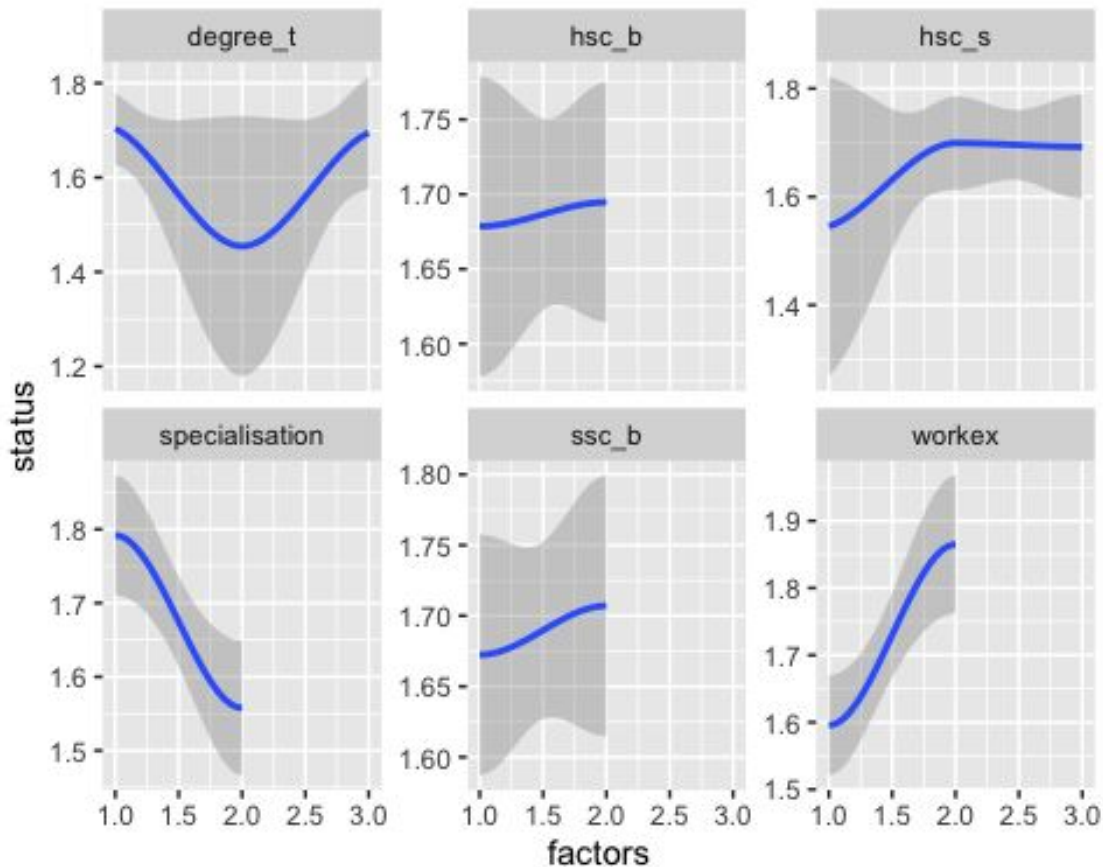
For better visualization of relationship with factor data, we are converting them to numerical

```
numPlacement<-placement
indx <- sapply(placement, is.factor)
numPlacement[indx] <- lapply(placement[indx], function(x) as.numeric(x))
```

Relationship with salary: Making use of the converted data to plot and show the relationship with salary

```
numPlacement%>%gather(numPlacement,factors,c("ssc_b","hsc_b",
"hsc_s","degree_t","workex","specialisation"))%>%
  ggplot(aes(x=factors,y=status))+geom_smooth()+ facet_wrap(~ numPlacement,
scales = "free_y")

## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



Linear Regression Model to predict the salary:

```
status_trainmodel_lm<-numPlacement[1:200,]
status_testmodel_lm<-numPlacement[201:215,]
status_trainmodel_lm_backup<-status_trainmodel_lm
status_trainmodel_lm<-lm(status~.,status_trainmodel_lm_backup)
status_pred<-predict(status_trainmodel_lm,status_testmodel_lm)
cor(status_testmodel_lm$status,status_pred)
```

```
## [1] 0.9480045
```

Comparing the status predicted with the test data: Rounding off the predicted value:

```
status_pred<-round(status_pred,digits = 0)
comparaison<-cbind(status_testmodel_lm$status,status_pred)
status_trainmodel_lm

##
## Call:
## lm(formula = status ~ ., data = status_trainmodel_lm_backup)
```

```
##
## Coefficients:
## (Intercept)          gender          ssc_p          ssc_b
hsc_p
##      8.010e-01      -9.249e-03      9.720e-03      9.400e-03
3.266e-03
##      hsc_b          hsc_s          degree_p          degree_t
workex
##      -3.096e-04      5.450e-03      8.357e-03      -5.145e-02
5.478e-02
##      etest_p  specialisation          mba_p          salary
##      -2.217e-03      7.392e-03      -1.286e-02      2.027e-06

comparaison

##      status_pred
## 201 2          2
## 202 1          1
## 203 2          2
## 204 2          2
## 205 2          2
## 206 2          2
## 207 1          1
## 208 2          2
## 209 1          1
## 210 2          2
## 211 2          2
## 212 2          2
## 213 2          2
## 214 2          2
## 215 1          1
```

Labeling the status values to concerning labels:

```
status_pred[status_pred==2]<-"Placed"
status_pred[status_pred==1]<-"Not Placed"
status_testmodel_lm$status[status_testmodel_lm$status==1]<-"Not Placed"
status_testmodel_lm$status[status_testmodel_lm$status==2]<-"Placed"
comparaison<-cbind(status_testmodel_lm$status,status_pred)
status_trainmodel_lm

##
## Call:
## lm(formula = status ~ ., data = status_trainmodel_lm_backup)
##
## Coefficients:
## (Intercept)          gender          ssc_p          ssc_b
hsc_p
##      8.010e-01      -9.249e-03      9.720e-03      9.400e-03
```

```

3.266e-03
##          hsc_b          hsc_s          degree_p          degree_t
workex
##      -3.096e-04      5.450e-03      8.357e-03      -5.145e-02
5.478e-02
##          etest_p  specialisation          mba_p          salary
##      -2.217e-03      7.392e-03      -1.286e-02      2.027e-06

comparaison

##          status_pred
## 201 "Placed"      "Placed"
## 202 "Not Placed" "Not Placed"
## 203 "Placed"      "Placed"
## 204 "Placed"      "Placed"
## 205 "Placed"      "Placed"
## 206 "Placed"      "Placed"
## 207 "Not Placed" "Not Placed"
## 208 "Placed"      "Placed"
## 209 "Not Placed" "Not Placed"
## 210 "Placed"      "Placed"
## 211 "Placed"      "Placed"
## 212 "Placed"      "Placed"
## 213 "Placed"      "Placed"
## 214 "Placed"      "Placed"
## 215 "Not Placed" "Not Placed"

```

Predict the salary using the Linear Regression Method:

```

salary_trainmodel_lm<-placement[1:200,]
salary_testmodel_lm<-placement[201:215,]
salary_trainmodel_lm_backup<-salary_trainmodel_lm
salary_trainmodel_lm<-lm(salary~.,salary_trainmodel_lm_backup)
salary_pred<-predict(salary_trainmodel_lm,salary_testmodel_lm)
cor(salary_testmodel_lm$salary,salary_pred)

## [1] 0.950907

```

Comparing the salary predicted with the test values:

```

comparaison<-cbind(salary_testmodel_lm$salary,salary_pred)
comparaison

##          salary_pred
## 201 300000 284034.510
## 202      0 -16464.768
## 203 240000 295375.737

```

```
## 204 260000 292923.099
## 205 210000 296408.474
## 206 250000 284792.899
## 207      0  13395.782
## 208 300000 334151.869
## 209      0 -12407.543
## 210 216000 291125.733
## 211 400000 316476.131
## 212 275000 298198.751
## 213 295000 311757.438
## 214 204000 265419.231
## 215      0 -5348.827
```

Predict using knn classification and compare using Cross Table:

Status prediction:

```
normalize <- function(x) {return ((x - min(x)) / ((max(x) - min(x))))}

numPlacement_n <- as.data.frame(lapply(numPlacement[,1:12], normalize))

status_train_Knn<-numPlacement_n[1:200,]
status_test_Knn<-numPlacement_n[201:215,]

status_train_labels<-placement[1:200,13]
status_test_labels<-placement[201:215,13]

status_pred_knn<-knn(train = status_train_Knn, test = status_test_Knn,
cl=status_train_labels, k=15)
```

```
CrossTable(x=status_test_labels,y=status_pred_knn,prop.chisq=FALSE)
```

```
##
##
##      Cell Contents
## |-----|
## |                      N |
## |      N / Row Total |
## |      N / Col Total |
## |      N / Table Total |
## |-----|
##
##
## Total Observations in Table:  15
##
##
##                | status_pred_knn
```



```
## status_test_labels | Not Placed | Placed | Row Total |
## -----|-----|-----|-----|
## Not Placed | 3 | 1 | 4 |
## | 0.750 | 0.250 | 0.267 |
## | 0.750 | 0.091 | |
## | 0.200 | 0.067 | |
## -----|-----|-----|
## Placed | 1 | 10 | 11 |
## | 0.091 | 0.909 | 0.733 |
## | 0.250 | 0.909 | |
## | 0.067 | 0.667 | |
## -----|-----|-----|
## Column Total | 4 | 11 | 15 |
## | 0.267 | 0.733 | |
## -----|-----|-----|
##
##
```

Salary prediction:

```
numPlacement_salary_n <- as.data.frame(lapply(numPlacement[,1:13],
normalize))

salary_train_Knn<-numPlacement_salary_n[1:200,]
salary_test_Knn<-numPlacement_salary_n[201:215,]

salary_train_labels<-placement[1:200,14]
salary_test_labels<-placement[201:215,14]

salary_pred_knn<-knn(train = salary_train_Knn, test = salary_test_Knn,
cl=salary_train_labels, k=15)

CrossTable(x=salary_test_labels,y=salary_pred_knn,prop.chisq=FALSE)
```

```
##
##
## Cell Contents
## |-----|
## | N |
## | N / Row Total |
## | N / Col Total |
## | N / Table Total |
## |-----|
##
##
## Total Observations in Table: 15
##
##
```

```

##                                     | salary_pred_knn
## salary_test_labels |           0 |      240000 |      250000 |      265000 |
3e+05 |      5e+05 | Row Total |
##
-----|-----|-----|-----|-----|-----
--|-----|-----|
##           0 |           4 |           0 |           0 |           0 |
0 |           0 |           4 |
##           |           1.000 |           0.000 |           0.000 |           0.000 |
0.000 |           0.000 |           0.267 |
##           |           0.800 |           0.000 |           0.000 |           0.000 |
0.000 |           0.000 |           |
##           |           0.267 |           0.000 |           0.000 |           0.000 |
0.000 |           0.000 |           |
##
-----|-----|-----|-----|-----|-----
--|-----|-----|
##      204000 |           0 |           0 |           0 |           1 |
0 |           0 |           1 |
##           |           0.000 |           0.000 |           0.000 |           1.000 |
0.000 |           0.000 |           0.067 |
##           |           0.000 |           0.000 |           0.000 |           0.500 |
0.000 |           0.000 |           |
##           |           0.000 |           0.000 |           0.000 |           0.067 |
0.000 |           0.000 |           |
##
-----|-----|-----|-----|-----|-----
--|-----|-----|
##      210000 |           0 |           0 |           0 |           0 |
1 |           0 |           1 |
##           |           0.000 |           0.000 |           0.000 |           0.000 |
1.000 |           0.000 |           0.067 |
##           |           0.000 |           0.000 |           0.000 |           0.000 |
0.500 |           0.000 |           |
##           |           0.000 |           0.000 |           0.000 |           0.000 |
0.067 |           0.000 |           |
##
-----|-----|-----|-----|-----|-----
--|-----|-----|
##      216000 |           0 |           0 |           0 |           0 |
1 |           0 |           1 |
##           |           0.000 |           0.000 |           0.000 |           0.000 |
1.000 |           0.000 |           0.067 |
##           |           0.000 |           0.000 |           0.000 |           0.000 |
0.500 |           0.000 |           |
##           |           0.000 |           0.000 |           0.000 |           0.000 |
0.067 |           0.000 |           |
##

```

----- ----- ----- ----- ----- -----						
-- ----- -----						
##	240000	1	0	0	0	
0	0	1				
##		1.000	0.000	0.000	0.000	
0.000	0.000	0.067				
##		0.200	0.000	0.000	0.000	
0.000	0.000					
##		0.067	0.000	0.000	0.000	
0.000	0.000					
##						
----- ----- ----- ----- ----- -----						
-- ----- -----						
##	250000	0	1	0	0	
0	0	1				
##		0.000	1.000	0.000	0.000	
0.000	0.000	0.067				
##		0.000	0.250	0.000	0.000	
0.000	0.000					
##		0.000	0.067	0.000	0.000	
0.000	0.000					
##						
----- ----- ----- ----- ----- -----						
-- ----- -----						
##	260000	0	0	0	1	
0	0	1				
##		0.000	0.000	0.000	1.000	
0.000	0.000	0.067				
##		0.000	0.000	0.000	0.500	
0.000	0.000					
##		0.000	0.000	0.000	0.067	
0.000	0.000					
##						
----- ----- ----- ----- ----- -----						
-- ----- -----						
##	275000	0	0	0	0	
0	1	1				
##		0.000	0.000	0.000	0.000	
0.000	1.000	0.067				
##		0.000	0.000	0.000	0.000	
0.000	1.000					
##		0.000	0.000	0.000	0.000	
0.000	0.067					
##						
----- ----- ----- ----- ----- -----						
-- ----- -----						
##	295000	0	1	0	0	
0	0	1				

```

##          | 0.000 | 1.000 | 0.000 | 0.000 |
0.000 | 0.000 | 0.067 |
##          | 0.000 | 0.250 | 0.000 | 0.000 |
0.000 | 0.000 |
##          | 0.000 | 0.067 | 0.000 | 0.000 |
0.000 | 0.000 |
##
-----|-----|-----|-----|-----|-----
--|-----|-----|
##          3e+05 | 0 | 1 | 1 | 0 |
0 | 0 | 2 |
##          | 0.000 | 0.500 | 0.500 | 0.000 |
0.000 | 0.000 | 0.133 |
##          | 0.000 | 0.250 | 1.000 | 0.000 |
0.000 | 0.000 |
##          | 0.000 | 0.067 | 0.067 | 0.000 |
0.000 | 0.000 |
##
-----|-----|-----|-----|-----|-----
--|-----|-----|
##          4e+05 | 0 | 1 | 0 | 0 |
0 | 0 | 1 |
##          | 0.000 | 1.000 | 0.000 | 0.000 |
0.000 | 0.000 | 0.067 |
##          | 0.000 | 0.250 | 0.000 | 0.000 |
0.000 | 0.000 |
##          | 0.000 | 0.067 | 0.000 | 0.000 |
0.000 | 0.000 |
##
-----|-----|-----|-----|-----|-----
--|-----|-----|
##      Column Total | 5 | 4 | 1 | 2 |
2 | 1 | 15 |
##          | 0.333 | 0.267 | 0.067 | 0.133 |
0.133 | 0.067 |
##
-----|-----|-----|-----|-----|-----
--|-----|-----|
##
##

```

Naive_bayes algorithm

Naive bayes can only be used to predict Categorical data and not continuous data. So, we can only predict Status using it.

status Prediction:

```
status_NB<-naiveBayes(as.factor(status)~.,data=placement)
status_prediction_NB<-predict(status_NB,placement[201:215,])
```

```
CrossTable(x=status_test_labels,y=status_prediction_NB,prop.chisq=FALSE)
```

```
##
##
##      Cell Contents
## |-----|
## |              N |
## |      N / Row Total |
## |      N / Col Total |
## |      N / Table Total |
## |-----|
##
##
## Total Observations in Table:  15
##
##
##      status_prediction_NB
## status_test_labels | Not Placed |      Placed | Row Total |
## -----|-----|-----|-----|
##      Not Placed |      4 |      0 |      4 |
##      |      1.000 |      0.000 |      0.267 |
##      |      0.333 |      0.000 |
##      |      0.267 |      0.000 |
## -----|-----|-----|
##      Placed |      8 |      3 |      11 |
##      |      0.727 |      0.273 |      0.733 |
##      |      0.667 |      1.000 |
##      |      0.533 |      0.200 |
## -----|-----|-----|
##      Column Total |      12 |      3 |      15 |
##      |      0.800 |      0.200 |
## -----|-----|-----|
##
##
```

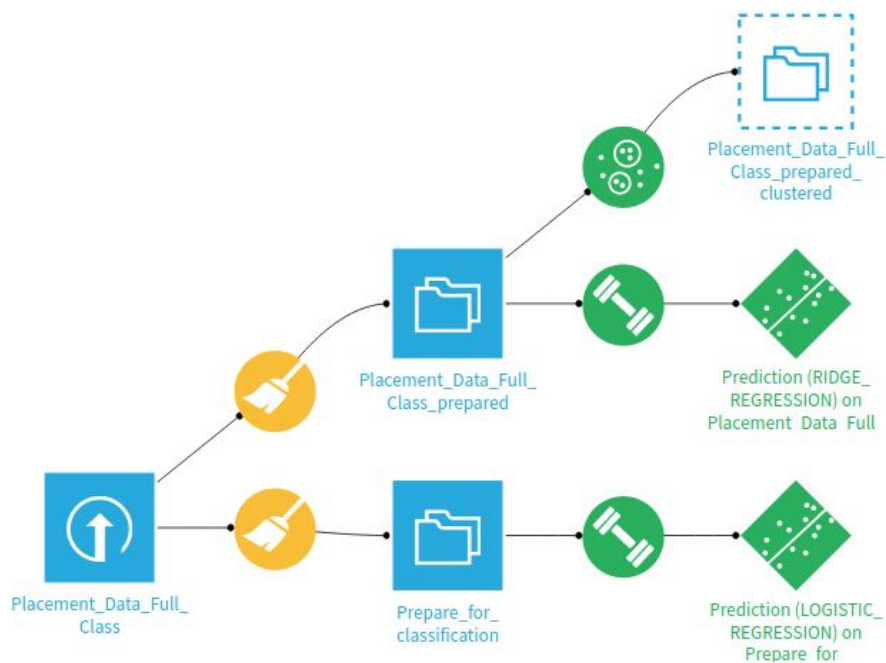
Dataiku:

The dataset was processed through different predictive models to determine the two different values of the dataset.

- Predict if the student was Placed or Not
- If the student was placed, what was the salary?


Flow:

Dataiku's intuitive interface makes it easier to describe the overall processing of data. You will find below, the overview of the data models designed and predictive analytics handled.



Regression to predict Salary:

Ridge (L2) Regression




This chart depicts a line chart comparing the predicted and actual values of salary for the test data

Error Distribution

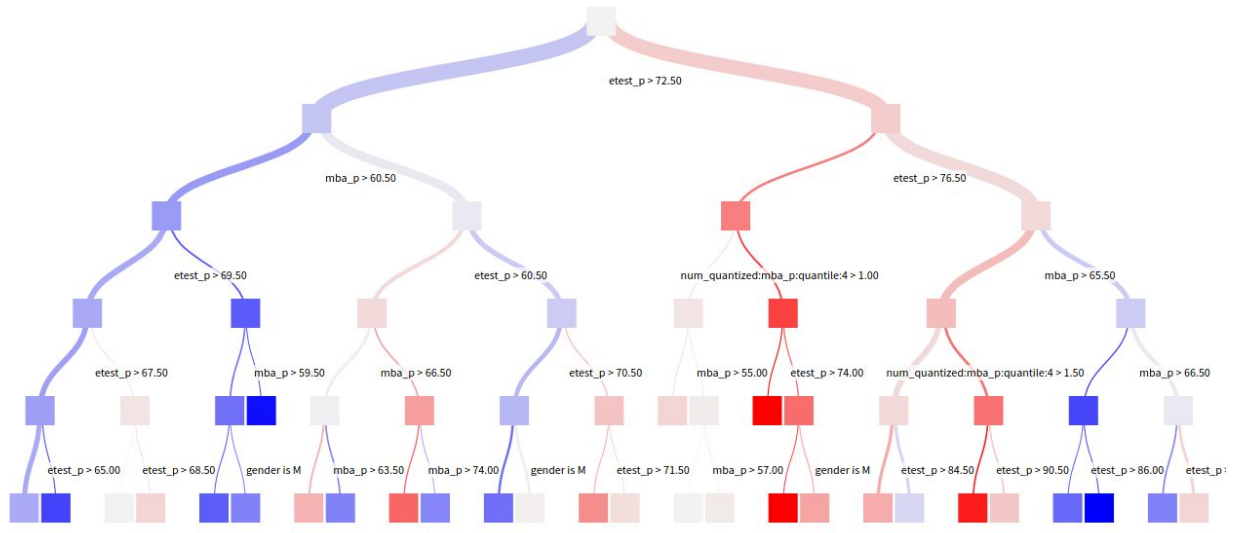
Min. (raw)	Min.	25 th perc.	Median	75 th perc.	90 th perc.	Max.	Max. (raw)
-78438	-75531	-16894	-2663.9	19450	61718	91059	1.0302e+5
Average		281.75	Standard deviation		42097		

The errors (difference between predicted and actual values) should be centered around zero, and the distribution should be "narrow", i.e the spread of the error should be limited. More generally, the errors should be "normally" distributed around zero (the curve should look like a bell).

To reduce the effect of possible spurious outliers, error distribution is winsorized (clipped) at the 2nd and 98th percentiles.




Decision Tree to predict Salary:



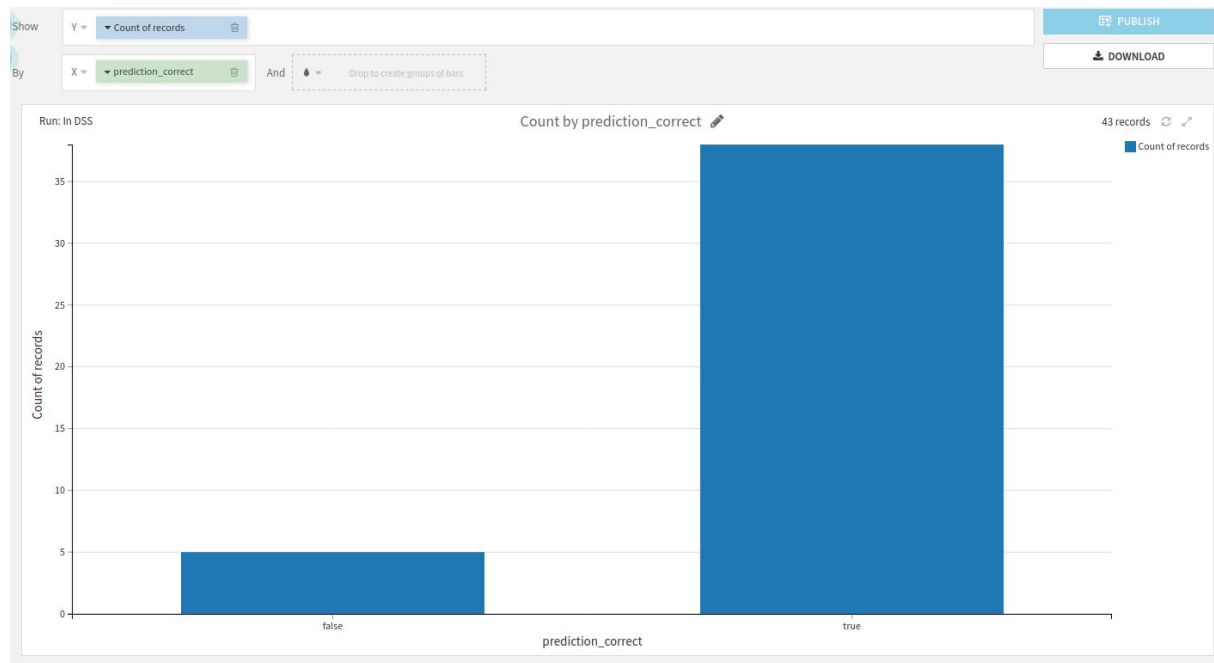
The Decision Tree was applied over the same dataset that was prepared for L2 regression.

Important columns that impact the prediction



Classification to predict Placed or not:


Logistic Regression



Threshold-dependent (current threshold = 0.8750)	
Accuracy	0.8837
Proportion of correct predictions (positive and negative) in the test set	
Precision	0.7692
Proportion of positive predictions that were indeed positive (in the test set)	
Recall	0.8333
Proportion of actual positive values found by the classifier	
F1 Score	0.8000
Harmonic mean between Precision and Recall	
Hamming loss	0.1163
Fraction of labels that are incorrectly predicted (the lower the better)	
Matthews Correlation Coefficient	0.7194
Correlation coefficient between actual and predicted values. +1 = perfect, 0 = no correlation, -1 = perfect anti-correlation	
Threshold-independent	
Log loss	0.5272
Error metric that takes into account the predicted probabilities (the lower the better)	
ROC - AUC Score	0.9382
Area under the ROC; from 0.5 (random model) to 1 (perfect model)	
Calibration loss	0.2303
Average distance between calibration curve and diagonal. From 0 (perfectly calibrated) up to 0.5.	

The Classification model was targeted to predict that “status” column of the dataset. From the model we can see that accuracy was very high and there is a distinct possibility of the data being overfitted


Clustering data to extract further info:



View features : **ALL** NUMERIC SINGLE CATEGORICAL



Click on a feature or a cluster to sort the heatmap by significance.



Here we can see that,

- Cluster_0 : Gender being Male is a significant factor for the datapoint
- Cluster_1 : The values of scores from ssc, hsc, mba and etest all are equally important. We also see that data points with Gender being *Female* is more likely to fall under this group
- Cluster_2 : There doesn't seem to be any significant associativity with any of the columns for this cluster

These clusters can be divided up into different datasets and can be used for cleaner and more accurate analysis of the patterns within them.